

Appendix

April 20, 2018

```
In [4]: import time
import numpy as np
import pandas as pd
from scipy import stats
pd.options.display.float_format = '{:,.3f}'.format

# %matplotlib notebook
import seaborn as sns
sns.set(style="ticks", color_codes=True)
import matplotlib.pyplot as plt

from sklearn.decomposition import PCA, NMF
from sklearn.manifold import TSNE
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error

from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.svm import SVR
from sklearn.ensemble import BaggingRegressor, ExtraTreesRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neural_network import MLPRegressor

class Data:
    #storage class for data
    def __init__(self, filename, subset_count=None):
        x = pd.read_csv(filename, index_col='id')
        if subset_count is not None: # Nice to use for testing
            permut = np.random.permutation(x.shape[0])[:subset_count]
            x = x.iloc[permut, :]

        self.df_y = x['loss']
        y = x['loss'].values
        del x['loss']

        convertedX = pd.get_dummies(x, drop_first=True)
```

```

X = convertedX.values
self.X = X
self.y = y

self.df_X = convertedX
self.df_X_test = None

def get_split(self, test_size=0.20, pca_components=None, nmf_components=None):
    X_train, X_test, y_train, y_test = train_test_split(self.X, self.y,
                                                         test_size=test_size)

    if pca_components is not None:
        pca = PCA(n_components=pca_components)
        pca.fit(X_train)
        X_train = pca.transform(X_train)
        X_test = pca.transform(X_test)
    elif nmf_components:
        nmf = NMF(n_components=nmf_components)
        X_train = nmf.fit_transform(X_train)
        X_test = nmf.transform(X_test)
    return X_train, X_test, y_train, y_test

def remove_columns(self, columns):
    columns = tuple(columns)
    for col in self.df_X.columns:
        if col.startswith(columns):
            del self.df_X[col]
    self.X = self.df_X.values

def read_test_data(self, filename):
    X_test = pd.read_csv(filename, index_col='id')
    X_test = pd.get_dummies(X_test, drop_first=True)

    additional_columns = set(X_test.columns) - set(self.df_X.columns)
    X_test = X_test.drop(columns=additional_columns)

    missing_columns = set(self.df_X.columns) - set(X_test.columns)
    for col in missing_columns:
        X_test[col] = 0

    self.df_X_test = X_test
    return X_test.values

def evaluate(name, estimator, X_train, X_test, y_train, y_test):
    t_0 = time.time()
    print(f'{name}:')
    estimator.fit(X_train, y_train)

    t_1 = time.time()

```

```

print(f'\tTime elapsed for model construction {t_1 - t_0:.3f} sec')
y_test_predict = estimator.predict(X_test)
error_test = mean_absolute_error(y_test, y_test_predict)
error_train = mean_absolute_error(y_train, estimator.predict(X_train))
print(f'\tTime elapsed for prediction {time.time() - t_1:.3f} sec')
print(f'\tTest error: {error_test:.3f}')
print(f'\tTrain error: {error_train:.3f}')
return error_test

```

```

In [ ]: # First overall test
data = Data("train.csv")
data.remove_columns(['cont9', 'cont12', 'cat2', 'cat3', 'cat4',
                    'cat5', 'cat6', 'cat7', 'cat8', 'cat86'])
X_train, X_test, y_train, y_test = data.get_split()

ESTIMATORS = {
    # Linear
    "LinearRegression": LinearRegression(n_jobs=-1),
    "Ridge": Ridge(),
    "Lasso": Lasso(),
    "ElasticNet": ElasticNet(),
    # Non-linear
    "BaggingRegressor": BaggingRegressor(n_jobs=-1),
    "ExtraTreesRegressor": ExtraTreesRegressor(n_jobs=-1),
    "RandomForestRegressor": RandomForestRegressor(n_jobs=-1),
    "GradientBoostingRegressor": GradientBoostingRegressor(loss='huber'),
    "MLP": MLPRegressor(),
    "KNeighborsRegressor": KNeighborsRegressor(n_jobs=-1),
    "SVR": SVR(),
}

for name, estimator in ESTIMATORS.items():
    evaluate(name, estimator, X_train, X_test, y_train, y_test)

```

```

LinearRegression:
    Time elapsed for model construction 23.044 sec
    Time elapsed for prediction 0.478 sec
    Test error: 8798305800.629
    Train error: 1480.444

```

```

Ridge:
    Time elapsed for model construction 5.009 sec
    Time elapsed for prediction 0.391 sec
    Test error: 1477.830
    Train error: 1481.719

```

```

Lasso:
    Time elapsed for model construction 55.786 sec
    Time elapsed for prediction 0.389 sec
    Test error: 1481.435

```

```

        Train error: 1493.583
ElasticNet:
    Time elapsed for model construction 6.423 sec
    Time elapsed for prediction 0.389 sec
    Test error: 1745.976
    Train error: 1764.948
BaggingRegressor:
    Time elapsed for model construction 136.550 sec
    Time elapsed for prediction 34.328 sec
    Test error: 1519.984
    Train error: 605.857
ExtraTreesRegressor:
    Time elapsed for model construction 279.494 sec
    Time elapsed for prediction 1.537 sec
    Test error: 1535.685
    Train error: 0.072
RandomForestRegressor:
    Time elapsed for model construction 139.742 sec
    Time elapsed for prediction 1.370 sec
    Test error: 1519.550
    Train error: 602.167
GradientBoostingRegressor:
    Time elapsed for model construction 822.191 sec
    Time elapsed for prediction 1.641 sec
    Test error: 1444.904
    Train error: 1453.168
MLP:
    Time elapsed for model construction 2397.275 sec
    Time elapsed for prediction 2.283 sec
    Test error: 1410.916
    Train error: 1392.052
KNeighborsRegressor:
    Time elapsed for model construction 119.529 sec
    Time elapsed for prediction 17134.402 sec
    Test error: 1672.010
    Train error: 1364.940
SVR:
    Time elapsed for model construction 20007.080 sec
    Time elapsed for prediction 10160.075 sec
    Test error: 1755.613
    Train error: 1724.410

In [ ]: ESTIMATORS = {
        # Linear
        "LinearRegression": LinearRegression(n_jobs=-1),
        "Ridge": Ridge(),
        "Lasso": Lasso(),

```

```

    "ElasticNet": ElasticNet(),
    # Non-linear
    "BaggingRegressor": BaggingRegressor(n_jobs=-1),
    "ExtraTreesRegressor": ExtraTreesRegressor(n_jobs=-1),
    "RandomForestRegressor": RandomForestRegressor(n_jobs=-1),
    "GradientBoostingRegressor": GradientBoostingRegressor(loss='huber'),
    "MLP": MLPRegressor(),
}

# Test PCA and NMF
data = Data("train.csv")

X_train, X_test, y_train, y_test = data.get_split(pca_components=120)
print(f'PCA with 120 components')
for name, estimator in ESTIMATORS.items():
    evaluate(name, estimator, X_train, X_test, y_train, y_test)

X_train, X_test, y_train, y_test = data.get_split(nmf_components=90)
print(f'NMF with 90 components')
for name, estimator in ESTIMATORS.items():
    evaluate(name, estimator, X_train, X_test, y_train, y_test)

```

PCA with 120 components

LinearRegression:

```

Time elapsed for model construction 1.222 sec
Time elapsed for prediction 0.052 sec
Test error: 1336.491
Train error: 1326.450

```

Ridge:

```

Time elapsed for model construction 0.375 sec
Time elapsed for prediction 0.047 sec
Test error: 1336.488
Train error: 1326.447

```

Lasso:

```

Time elapsed for model construction 0.735 sec
Time elapsed for prediction 0.050 sec
Test error: 1335.520
Train error: 1325.040

```

ElasticNet:

```

Time elapsed for model construction 0.688 sec
Time elapsed for prediction 0.059 sec
Test error: 1503.928
Train error: 1487.781

```

BaggingRegressor:

```

Time elapsed for model construction 159.103 sec
Time elapsed for prediction 7.503 sec
Test error: 1385.234
Train error: 551.043

```

ExtraTreesRegressor:
Time elapsed for model construction 37.692 sec
Time elapsed for prediction 0.782 sec
Test error: 1379.454
Train error: 0.001

RandomForestRegressor:
Time elapsed for model construction 151.161 sec
Time elapsed for prediction 0.754 sec
Test error: 1390.771
Train error: 552.164

GradientBoostingRegressor:
Time elapsed for model construction 333.502 sec
Time elapsed for prediction 0.793 sec
Test error: 1317.153
Train error: 1290.940

MLP:
Time elapsed for model construction 327.754 sec
Time elapsed for prediction 1.168 sec
Test error: 1224.814
Train error: 1195.139

NMF with 90 components

LinearRegression:
Time elapsed for model construction 0.858 sec
Time elapsed for prediction 0.047 sec
Test error: 1372.249
Train error: 1369.867

Ridge:
Time elapsed for model construction 0.288 sec
Time elapsed for prediction 0.031 sec
Test error: 1367.990
Train error: 1366.349

Lasso:
Time elapsed for model construction 1.047 sec
Time elapsed for prediction 0.047 sec
Test error: 1380.740
Train error: 1380.795

ElasticNet:
Time elapsed for model construction 0.509 sec
Time elapsed for prediction 0.047 sec
Test error: 1969.046
Train error: 1961.323

BaggingRegressor:
Time elapsed for model construction 63.872 sec
Time elapsed for prediction 6.347 sec
Test error: 1350.245
Train error: 536.754

ExtraTreesRegressor:

```

Time elapsed for model construction 28.637 sec
Time elapsed for prediction 0.648 sec
Test error: 1359.542
Train error: 0.009
RandomForestRegressor:
Time elapsed for model construction 59.382 sec
Time elapsed for prediction 0.608 sec
Test error: 1344.332
Train error: 537.979
GradientBoostingRegressor:
Time elapsed for model construction 150.457 sec
Time elapsed for prediction 0.901 sec
Test error: 1263.213
Train error: 1253.464
MLP:
Time elapsed for model construction 1372.075 sec
Time elapsed for prediction 6.929 sec
Test error: 1363.668
Train error: 1361.900

In [ ]: data = Data("train.csv")
X_train, X_test, y_train, y_test = data.get_split()
for n_estimators in np.arange(200, 601, 50):
    estimator = GradientBoostingRegressor(loss='huber', alpha=0.5,
                                           n_estimators=n_estimators, max_depth=6,
                                           learning_rate=0.1, min_samples_leaf=10,
                                           min_samples_split=10)
    evaluate(f'GBR n_estimators={n_estimators}',
            estimator, X_train, X_test, y_train, y_test)

GBR n_estimators=200:
Time elapsed for model construction 9316.946 sec
Time elapsed for prediction 4.889 sec
Test error: 1163.757
Train error: 1086.791
GBR n_estimators=250:
Time elapsed for model construction 5754.714 sec
Time elapsed for prediction 4.712 sec
Test error: 1162.030
Train error: 1076.098
GBR n_estimators=300:
Time elapsed for model construction 6771.478 sec
Time elapsed for prediction 5.666 sec
Test error: 1160.265
Train error: 1066.413
GBR n_estimators=350:
Time elapsed for model construction 23563.061 sec

```

```

        Time elapsed for prediction 7.427 sec
        Test error: 1159.374
        Train error: 1057.450
GBR n_estimators=400:
    Time elapsed for model construction 8711.057 sec
    Time elapsed for prediction 7.002 sec
    Test error: 1158.555
    Train error: 1052.084
GBR n_estimators=450:
    Time elapsed for model construction 9692.810 sec
    Time elapsed for prediction 7.645 sec
    Test error: 1158.072
    Train error: 1046.006
GBR n_estimators=500:
    Time elapsed for model construction 10792.895 sec
    Time elapsed for prediction 8.363 sec
    Test error: 1158.072
    Train error: 1039.049
GBR n_estimators=550:
    Time elapsed for model construction 12230.542 sec
    Time elapsed for prediction 10.199 sec
    Test error: 1152.069
    Train error: 1031.342
GBR n_estimators=600:
    Time elapsed for model construction 14028.412 sec
    Time elapsed for prediction 14.199 sec
    Test error: 1142.065
    Train error: 1025.139

```

```

In [17]: # Reading in the final test data
        data = Data("train.csv")
        X_train, y_train = data.df_X, data.df_y
        X_test = data.read_test_data("test.csv")
        X_test = data.df_X_test

        estimator = GradientBoostingRegressor(loss='huber', alpha=0.5,
                                              n_estimators=600, max_depth=6,
                                              learning_rate=0.1, min_samples_leaf=10,
                                              min_samples_split=10)

        t_0 = time.time()
        estimator.fit(X_train, y_train)

        t_1 = time.time()
        print(f'Time elapsed for model construction {t_1 - t_0:.3f} sec')

        y_test_predict = estimator.predict(X_test)

```



```
X_test['loss'] = y_test_predict
X_test['loss'].to_csv("GBR-submission.csv", header=True)
```

Time elapsed for model construction 10289.486 sec

```
In [21]: data = Data("train.csv")
X_train, y_train = data.df_X, data.df_y
X_test = data.read_test_data("test.csv")
X_test = data.df_X_test

nmf = NMF(n_components=90)
X_train_trans = nmf.fit_transform(X_train)
X_test_trans = nmf.transform(X_test)

gbr = GradientBoostingRegressor(loss='huber', alpha=0.5,
                                n_estimators=600, max_depth=6,
                                learning_rate=0.1, min_samples_leaf=10,
                                min_samples_split=10)

t_0 = time.time()
gbr.fit(X_train_trans, y_train)

t_1 = time.time()
print(f'Time elapsed for model construction {t_1 - t_0:.3f} sec')

y_test_predict = gbr.predict(X_test_trans)

X_test['loss'] = y_test_predict
X_test['loss'].to_csv("GBR-NMF-submission.csv", header=True)
```

Time elapsed for model construction 3228.545 sec

```
In [24]: X_train, X_test, y_train, y_test = data.get_split()
X_train = nmf.transform(X_train)
X_test = nmf.transform(X_test)

y_test_predict = gbr.predict(X_test)
y_train_predict = gbr.predict(X_train)


# Plot of predicted values versus the true values
plt.figure()
plt.scatter(y_test, y_test_predict)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r')
plt.ylabel('Predicted - loss')
plt.xlabel('Actual - loss')
```

```

# Plot of residuals
plt.figure()
plt.scatter(y_train_predict, stats.zscore(y_train_predict - y_train),
            c='b', s=40, alpha=0.5)
plt.scatter(y_test_predict, stats.zscore(y_test_predict - y_test),
            c='g', s=40)

plt.hlines(y=0, xmin=0, xmax=20)
plt.title('Residual Plot using training (blue) and test(green) data')
plt.ylabel('Residuals')
plt.xlabel('Loss')
plt.show()

```





```
In [20]: data = Data("train.csv")
X_train, y_train = data.df_X, data.df_y
X_test = data.read_test_data("test.csv")
X_test = data.df_X_test

nmf = NMF(n_components=90)
X_train_trans = nmf.fit_transform(X_train)
X_test_trans = nmf.transform(X_test)

y_train_trans = np.log1p(y_train)

gbr = GradientBoostingRegressor(loss='huber', alpha=0.5,
                                n_estimators=600, max_depth=6,
                                learning_rate=0.1, min_samples_leaf=10,
                                min_samples_split=10)

t_0 = time.time()
gbr.fit(X_train_trans, y_train_trans)

t_1 = time.time()
print(f'Time elapsed for model construction {t_1 - t_0:.3f} sec')

y_test_predict = np.expml(gbr.predict(X_test_trans))
```

```
X_test['loss'] = y_test_predict
X_test['loss'].to_csv("GBR-NMF-log-submission.csv", header=True)
```


Time elapsed for model construction 2696.148 sec

```
In [6]: X_train = pd.read_csv("train.csv", index_col='id')

# Count the number of items for each categorical attribute
cat_index = [x for x in X_train.columns if x.startswith('cat')]

n = 8
r = len(cat_index)//n + 1
plt.figure(figsize=(18, 30))
for i in range(r):
    for j in range(n if i != r-1
                    else ((len(cat_index)%n))
                    ):
        if (i*n + j + 1) > len(cat_index):
            break
        plot = plt.subplot(r, n, (i*n + j + 1))
        ax = sns.countplot(x=cat_index[i*n + j], data=X_train)
        plt.subplots_adjust(wspace=0.55, hspace=0.4, top=0.96)
        ax.set_ylabel('')
plt.suptitle('The counts for each categorical attribute')
plt.show()
```

The counts for each categorical attribute



```

In [8]: cat_index = [x for x in X_train.columns if x.startswith('cat')]

def letter_to_num(series):
    ret = []
    for c in series:
        if len(c) == 1:
            ret.append(ord(c.lower()) - 96)
        else:
            ret.append('')
            for s in c:
                ret[-1] += str(ord(s.lower()) - 96)
            ret[-1] = int(ret[-1])
    return ret

integer_categories = X_train[cat_index].apply(letter_to_num)

# spearman and kendall-tau mostly agree
corr = integer_categories.corr(method='kendall')

corr_matrix = corr.abs()
sorted_corr = (corr_matrix.where(np.triu(np.ones(corr_matrix.shape),
                                              k=1).astype(np.bool))
               .stack()
               .sort_values(ascending=False))
print('The most correlated categorical attributes are:')
sorted_corr[sorted_corr > 0.87]

```

The most correlated categorical attributes are:

```

Out[8]: cat7    cat89    0.999
        cat3    cat90    0.998
        cat8    cat102   0.991
        cat4    cat111   0.934
        cat2    cat9     0.932
        cat6    cat50    0.926
                cat114   0.925
        cat5    cat103   0.925
        cat86   cat99    0.888
dtype: float64

```

```

In [9]: cont_index = lambda : [x for x in X_train.columns if x.startswith('cont')]
plt.figure(figsize=(18, 10))

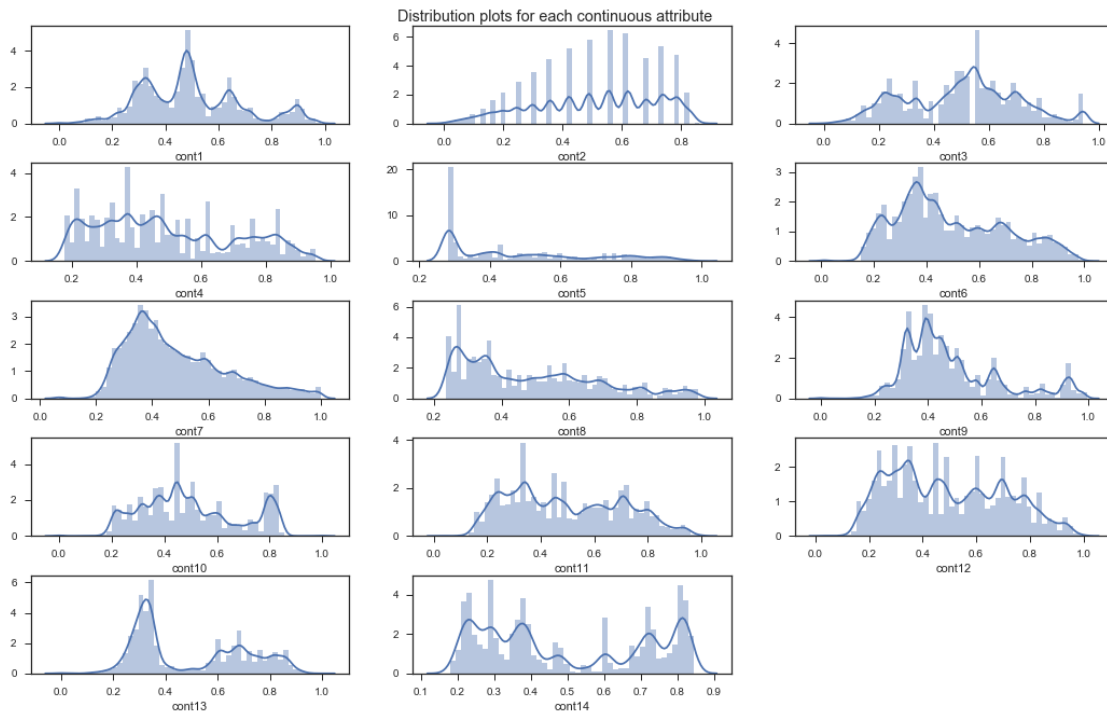
n = len(cont_index())

```

```

for i, col in enumerate(cont_index()):
    plot = plt.subplot(n//3+1, 3, i + 1)
    sns.distplot(X_train[col])
    plot.set_xlabel(f'{col}')
    plt.subplots_adjust(hspace=0.4, top=0.96)
plt.suptitle('Distribution plots for each continuous attribute')
plt.show()

```



```

In [ ]: #Keras code
import tensorflow, time
import os
import numpy
from keras.models import Sequential
from keras.layers import Dense, Activation, BatchNormalization
from keras.layers import Dropout
import keras.callbacks
from keras import optimizers
from keras import regularizers
from keras.constraints import maxnorm
from sklearn.decomposition import PCA
from sklearn.feature_selection import VarianceThreshold, SelectKBest, f_regression
from sklearn.model_selection import train_test_split

def TestEstimator(name, estimator, X_train, y_train, X_test, y_test, tf = False):

```

```

t=time.time()
if tf:
    y_train_ = numpy.log1p(y_train)
else:
    y_train_ = y_train

estimator.fit(X_train, y_train_)
trainTime = time.time()-t
t = time.time()

y_pred = estimator.predict(X_test)
if tf:
    y_pred = numpy.expm1(y_pred)
predictTime = time.time()-t

y_train_pred = estimator.predict(X_train)
if tf:
    y_train_pred = numpy.expm1(y_train_pred)

error = numpy.sum(numpy.abs(y_pred - y_test)) / len(y_test)
error_train = numpy.sum(numpy.abs(y_train_pred - y_train)) / len(y_train)
#print(time.ctime(),name,error,trainTime,predictTime)
return error_train, error #a difference here indicates divergence

EPOCHS = 100
class CustomMLPRegressor:
    """provide a similar interface to the sklearn regressors"""
    def LoadFromFile(self, filename):
        self.model = keras.models.load_model(filename)

    def __init__(self, nodes, earlyStopping, dropout):

        if type(nodes) is int:
            nodes = (nodes,)
        self.nodes = nodes
        self.dropout = dropout
        self.earlyStopping = earlyStopping

    def fit(self, X, y):
        #print (X,y)
        inputDim = X.shape[1]
        model = Sequential()
        self.model = model
        for i, nodeCount in enumerate(self.nodes):
            if i == 0:
                if len(self.nodes) == 1:
                    self.model.add(Dense(nodeCount,
                                           input_dim = inputDim,

```



```

                                activation='relu'))
        else:
            self.model.add(Dense(nodeCount,
                                   input_dim = inputDim,
                                   activation='relu',
                                   kernel_constraint = maxnorm(3)))
            if self.dropout:
                model.add(Dropout(self.dropout))
        else:
            if len(self.nodes) > i+1:
                self.model.add(Dense(nodeCount,
                                       activation='relu',
                                       kernel_constraint = maxnorm(3)))
            if self.dropout:
                model.add(Dropout(self.dropout))
            else:
                self.model.add(Dense(nodeCount, activation='relu'))
model.add(Dense(1))
opt = optimizers.Nadam()
model.compile(loss='mae', optimizer=opt)
#self.model = model

callback_list = []
filepath="weights-best.hdf5"
checkpoint = keras.callbacks.ModelCheckpoint(filepath,
                                              monitor='val_loss',
                                              verbose=0, save_best_only=True,
                                              mode='min')

callback_list.append(checkpoint)
earlyStopper = keras.callbacks.EarlyStopping(min_delta= self.earlyStopping,
                                              patience=5, verbose=0)
callback_list.append(earlyStopper)

history=model.fit(X.values.astype('float32'), y.astype('float32'),
                  epochs=EPOCHS, batch_size=1000,
                  verbose=2, validation_split=0.2, callbacks=callback_list)
bestModel = keras.models.load_model('weights-best.hdf5')
self.history = history
self.model =bestModel
return self
def predict(self, X):
    #print ('predicting')
    ret = self.model.predict(X.astype('float32'))
    #print (ret.shape)
    return ret.flatten()
def TestMLP(mlp):
    e_train, e_test = TestEstimator('mlp',mlp, X_train, y_train,
                                    X_test, y_test, tf=False)

```

```

print("mlp-%s-%.2f: %.2f %.2f"%(mlp.nodes,mlp.dropout,e_train, e_test))
try:
    os.rename('weights-best.hdf5',
              'weights-best-%.2f-%.2f-%s.hdf5'%(e_test,mlp.dropout,mlp.nodes))
except FileExistsError:
    pass
return e_train, e_test
#the following code was used to populate a results list.
# The lowest result for nodes/dropout pair was used to represent that
# point on the dataplot.
for i in range(2):
    nodes = 425+i*75
    for j in range(3):
        dropout =0.2+j*0.05
        mlp = CustomMLPRegressor((nodes,50,10),0.0,dropout)
        e_train, e = TestMLP(mlp)
        results.append([nodes,dropout,e_train, e])
        if e < bestE:
            bestE = e
            bestMLP = mlp

```